



Carnegie
Mellon
University

Sparse Training Data

Tutorial of Parameter Server

Mu Li

CSD@CMU & IDL@Baidu

muli@cs.cmu.edu

High-dimensional data are sparse

- ♦ Why high dimension?
 - ★ make the classifier's job easier
 - ★ linear method is often good enough
- ♦ Why sparse?
 - ★ easy to storage (only store non-zero entries)
 - ★ affordable computation cost
- ♦ Key difference to dense data when using
 - ★ require a lot of random read/write

Store and Computing

Compressed storage

	0	1	2
0		10	20
1	3		
2		87	
3	57		

♦ Sparse row-major:

offset = [0, 2, 3, 4, 5]

index = [1, 2, 0, 1, 0]

value = 10 20 3 87 57

♦ Sparse column-major:

offset = [0, 2, 4, 5]

index = [1, 3, 0, 2, 0]

value = 3 57 10 87 20

♦ Access $a(i,j)$ under row-major:

$k = \text{binary_search}(\text{index}[\text{offset}[i]], \text{index}[\text{offset}[i+1]], j)$
return valid(k) ? value(offset[i]+k) : 0

$$y = Ax$$

all offset, index, and value are read sequentially

♦ Sample C++ codes:

```
// matrix-vector multiplication y = A * x
void mat_vec(const V* x const, V* y) const {
    if (x == 0 || y == 0) return;
    for (size_t i = 0; i < rows(); ++i) { // i-th row
        for (size_t j = offset[i]; j < offset[i+1]; ++j)
            y[i] += x[index[j]] * value[j];
    }
}

// vector-matrix multiplication y = A * x
void vec_mat(const V* x const, V* y) const {
    if (x == 0 || y == 0) return;
    for (size_t i = 0; i < cols(); ++i) { // i-th column
        V x_i = x[i];
        for (size_t j = offset[i]; j < offset[i+1]; ++j)
            y[index[j]] += x_i * value[j];
    }
}
}
```

write y sequentially,
but read x in random

read x sequentially,
but write y in random

Numbers Everyone Should Know

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	100 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	10,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from network	10,000,000 ns
Read 1 MB sequentially from disk	30,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

10 times {

10 times {

slides by Jeff Dean



Cost of $y = Ax$

- ♦ The computation cost is $O(nnz(A))$
- ♦ The random access dominates the cost:
 $\approx \text{L2-cache-reference}(nnz(A))$
- ♦ In theory: process **1.4e8** nnz entries per second
- ♦ In reality: **8.4e7** nnz entries per second
 - ★ 4.3M x 17.4M sparse matrix
 - ★ mac mbp, Intel i7 2.3GHz cpu
 - ★ single thread

Real data

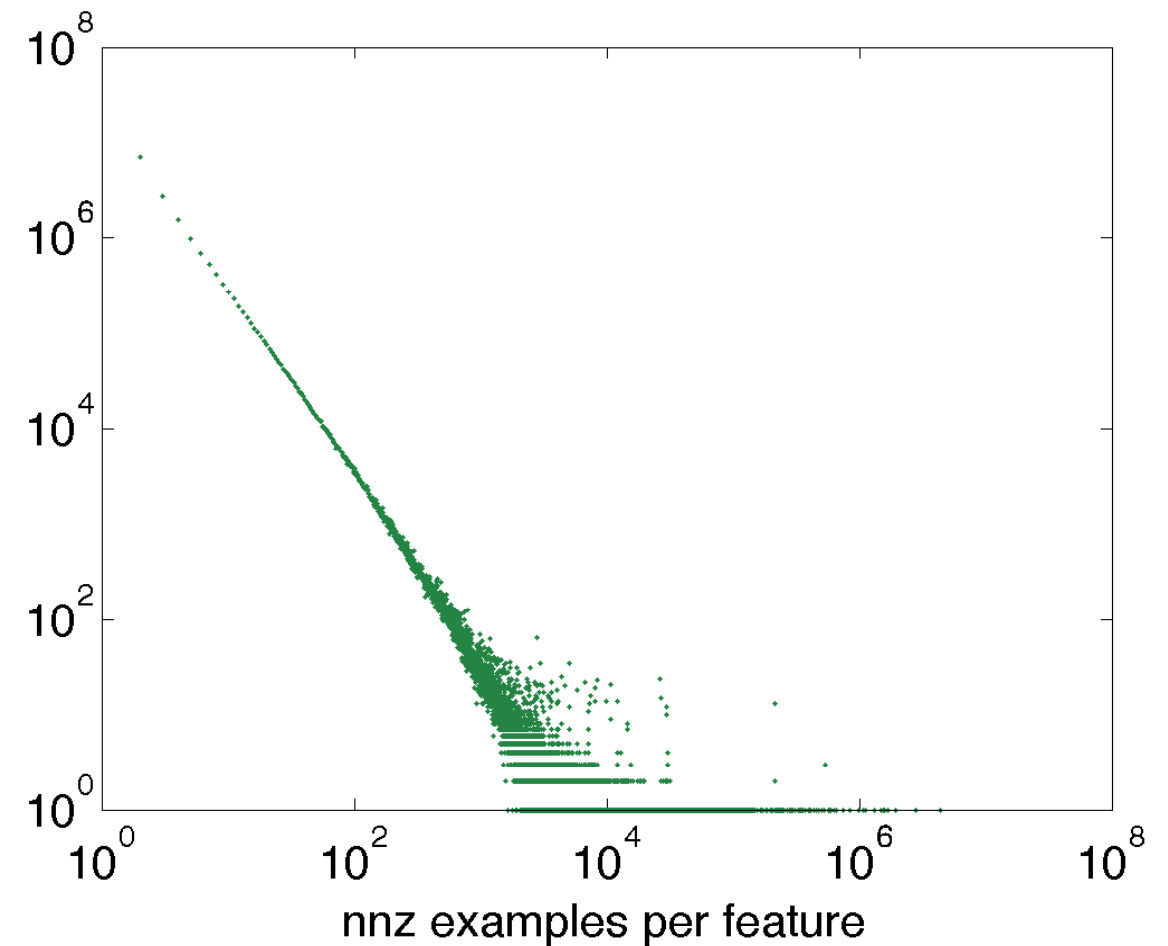

Product	Examples	Training data	Features per example
A	59.9B	2.00TB	54.9
B	7.6B	0.71TB	94.9
C	197.5B	15.54TB	77.7
D	129.1B	17.24TB	100.57

from sibyl

- ♦ ≈ 100 features per example is reasonable
- ★ \approx feature groups
- ★ for $y = Ax$, process $1e6$ examples per second
- ★ for linear method, 1000 cores, 100 billion examples, 100 iterations, finish in 3h in ideal

$$10^{11} \text{ examples} \times 100 \text{ iterations} / 1000 \text{ cores} / 10^6 = 1000 \text{ second}$$

Patterns of Sparsity

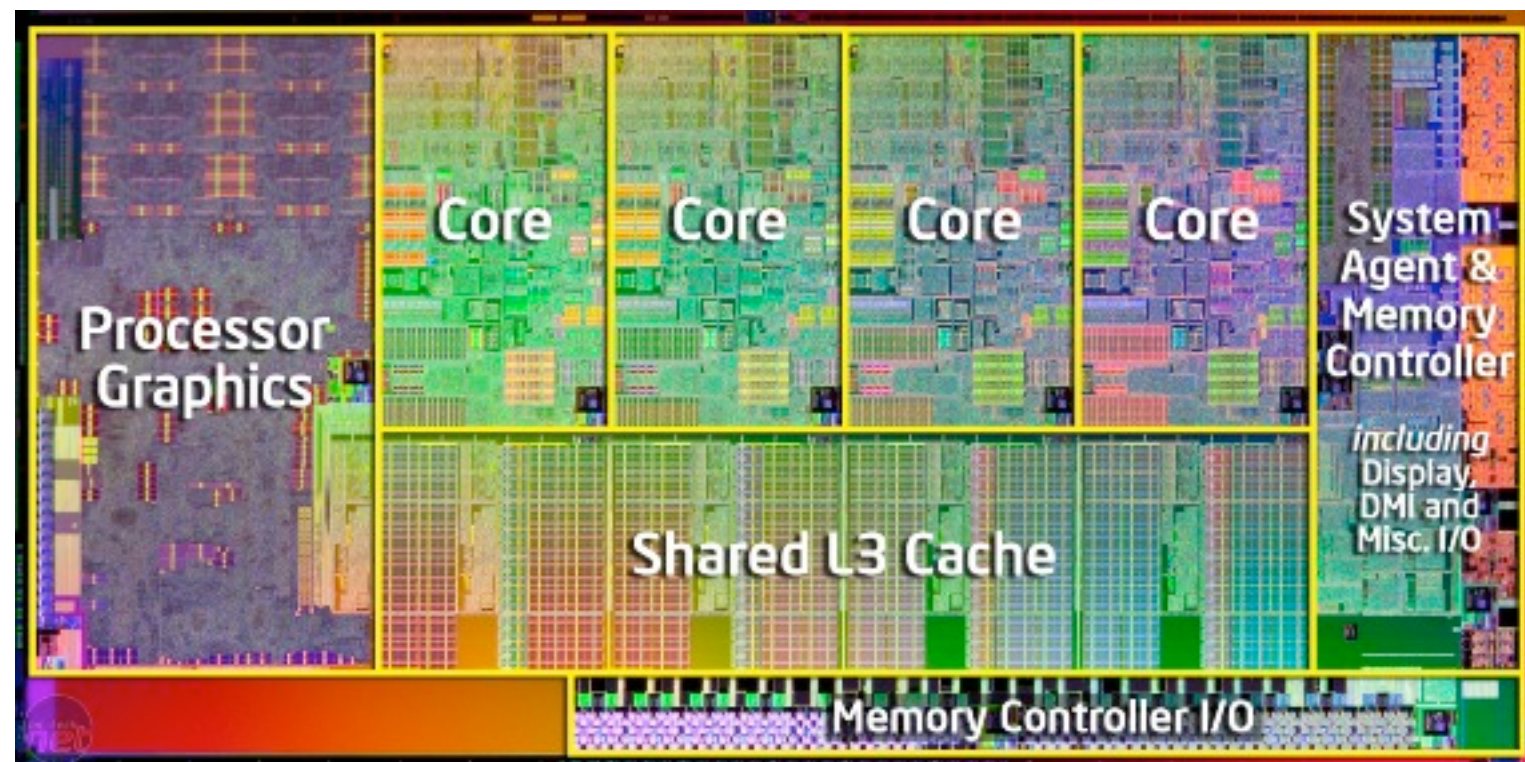


- ♦ non-zero entries are distributed irregularly on features
 - ★ imbalanced workload partition
 - ★ ill conditional number

Multi-thread Implementation

CPU


- ♦ Multiple cores (4-8)
- ♦ Multiple sockets (1-4)
- ♦ 2-4 GHz clock
- ♦ Memory interface 20-40GB/s
- ♦ Internal bandwidth $>100\text{GB/s}$



Benefits of multi-thread

fetch data from memory
~100 cycles

- ♦ Use more computation units
 - ★ float point units
- ♦ Hide the memory latency
 - ★ run something else when the data are not ready



Using ThreadPool

- ✦ A pool of threads, each one keeps fetching and executing unfinished tasks
- ✦ Create a pool with n threads: `ThreadPool pool(n)`
- ✦ Add a task into the pool: `pool.add(task)`
- ✦ Start executing: `pool.startWorkers()`



Multi-threaded $y = Ax$

- ✦ Assume row major
- ✦ Compute a segment of y

```
void rangeTimes(SizeR row_range, const V* const x, V* y) const;
```

- ✦ Divide y into several segments, each one is assigned to a thread

```
ThreadPool pool(num_threads);  
int num_tasks = rowMajor() ? num_threads * 10 : num_threads;  
for (int i = 0; i < num_tasks; ++i) {  
    pool.add([this, x, y, row_range, num_tasks, i]() {  
        rangeTimes(row_range.evenDivide(num_tasks, i), x, y);  
    });  
}  
pool.startWorkers();
```


c++11 lambda
functions

How about column-major?

- ✦ Equivalence to $x = A^T y$ for row-major A
- ✦ multi-threads concurrently write the same y
- ✦ Several possible solutions:
 - ★ convert into a row-major matrix first
 - ★ lock $y[i]$ (or a segment) before write it
 - ★ each thread writes only a segment of y

Experiments

- ♦ data: CTRa
- ♦ row major,
4.3M rows,
17.4M columns,
354M nnz entries
- ♦ MBP Pro,
Intel i7 2.3GHz,
4 cores,
8 hyper-threads



Coding Practice

- ♦ Implement $x = A^T y$
- ♦ You can reuse the codes at
 - ★ https://github.com/mli/mlss14_a
- ♦ CTRa in binary format is provided
 - ★ CTRa_X.index: 354700138 uint32
 - ★ CTRa_X.offset: 4349786 uint64
 - ★ CTRa_X.info: information
 - ★ 0-1 values, so ignore the value



Row major or column major

- ✦ No big difference for individual and whole access
 - ★ choose the one how data are stored
- ✦ Use row major when need read individual rows
 - ★ SGD, minibatch SGD, online learning
- ✦ Use column major when need read columns
 - ★ (block) Coordinate descent
- ✦ Converting cost $2 * \text{nnz}(A)$ random access

More

- ♦ Other operations? BLAS, LAPACK
 - ★ Timothy A. Davis, Direct Methods for Sparse Linear Systems, SIAM, 2006
- ♦ Existing packages:
 - ★ SuiteSparse, Eigen3...
 - ★ Use them as much as possible
 - ★ however, problem-specific optimizations may improve the performance a lot, we will see later

Eigen3

- ✦ Easy to install: all header files, just copy to a proper place
- ✦ Not easy to read: a lot of templates
- ✦ Good performance on dense data
- ✦ Somewhat convenient to use
- ✦ Jeff Dean is using it...

Bug 613 - Bug in internal::psqrt SSE implementation

Status: RESOLVED FIXED

Reported: 2013-06-13 17:54
UTC by **Jeff Dean**

Product: Eigen

Component: Core - general

Modified: 2013-06-14 09:52
UTC ([History](#))

add	<code>mat3 = mat1 + mat2;</code>	<code>mat3 += mat1;</code>
subtract	<code>mat3 = mat1 - mat2;</code>	<code>mat3 -= mat1;</code>
scalar product	<code>mat3 = mat1 * s1;</code> <code>mat3 = mat1 / s1;</code>	<code>mat3 *= s1;</code> <code>mat3 /= s1;</code>
matrix/vector products *	<code>col2 = mat1 * col1;</code> <code>row2 = row1 * mat1;</code> <code>mat3 = mat1 * mat2;</code>	<code>row1 *= mat1;</code> <code>mat3 *= mat1;</code>
transposition adjoint *	<code>mat1 = mat2.transpose();</code> <code>mat1 = mat2.adjoint();</code>	<code>mat1.transposeInPlace();</code> <code>mat1.adjointInPlace();</code>
dot product inner product *	<code>scalar = vec1.dot(vec2);</code> <code>scalar = col1.adjoint() * col2;</code> <code>scalar = (col1.adjoint() * col2).value();</code>	
outer product *	<code>mat = col1 * col2.transpose();</code>	